

Modular



# Mojo GPU Compilation



Weiwei Chen  
[weiwei.chen@modular.com](mailto:weiwei.chen@modular.com)

MLIR Workshop @  
LLVM Developers' Meeting 2025

# Agenda

---

- 01 Mojo in a glance 🔥
  - 02 Writing GPU code in Mojo 🖥️
  - 03 Mojo GPU Compilation
  - 04 Compiler Support for Kernel Debugging
  - 05 Conclusions
- 



# Mojo at a glance

Pythonic [systems programming](#) language

- CLI toolchain + VS Code support
- Strong Python and C interop
- Extensive generic programming\*, type system, and memory safety

Unlocking performance for heterogeneous compute

- Blazing fast
- Unified programming CPU + GPU
- Scalable GPU kernels for Nvidia, AMD, Apple ...

Best way to extend Python to CPUs and GPUs

Backbone for Modular's MAX Inference Engine  \*\*

```
def mandelbrot_kernel[  
    width: Int # SIMD Width  
](c: ComplexSIMD[float, width]) ->  
    SIMD[int, width]:  
  
        """A vectorized implementation of  
        the inner mandelbrot computation."""  
        z = ComplexSIMD[float, width](0, 0)  
  
        iters = SIMD[DType.index, width](0)  
        mask = SIMD[DType.bool, width](True)  
  
        for i in range(MAX_ITERS):  
            if not any(mask):  
                break  
  
            mask = z.squared_norm() <= 4  
            iters = mask.select(iters + 1, iters)  
            z = z.squared_add(c)  
        return iters
```

[Mojo !\[\]\(3e2231b1ad3ca8da8658228c00dd08e0\_img.jpg\) : A system programming language for heterogenous computing](#)  
LLVM Dev Meeting 2023


[Mojo : Programming language for all of AI](#)

[Mojo: A novel programming language for AI](#) ACAT 202

\* 10:30 Tue - *Building Modern Language Frontends with MLIR: Lessons from Mojo's Compile-Time Meta-Programming*

\*\* 10:30 Wed - *Modular MAX's JIT Graph Compiler*

# GPU Programming in Mojo



## GPU kernel in Mojo

Library imports

constants

GPU function

GPU tensor representation

GPU global offset

```
from math import ceildiv
from sys import has_accelerator
from gpu import global_idx
from gpu.host import DeviceContext
from layout import Layout, LayoutTensor

alias float_dtype = Dtype.float32
alias VECTOR_WIDTH = 10
alias BLOCK_SIZE = 5
alias layout = Layout.row_major(VECTOR_WIDTH)

fn vector_addition(
    lhs_tensor: LayoutTensor[float_dtype, layout, MutableAnyOrigin],
    rhs_tensor: LayoutTensor[float_dtype, layout, MutableAnyOrigin],
    out_tensor: LayoutTensor[float_dtype, layout, MutableAnyOrigin],
    size: Int,
):
    """The calculation to perform across the vector on the GPU."""
    var global_tid = global_idx.x
    if global_tid < UInt(size):
        out_tensor[global_tid] = lhs_tensor[global_tid] + rhs_tensor[global_tid]
```

## CPU host code in Mojo

static assertion

GPU DeviceContext

GPU device buffers

GPU tensors

compile and launch GPU kernel

device buffer to host

```
def main():
    constrained[has_accelerator(), "This example requires a supported GPU"]()

    # Get context for the attached GPU
    var ctx = DeviceContext()

    # Allocate data on the GPU address space
    var lhs_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)
    var rhs_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)
    var out_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)

    # Fill in values across the entire width
    _ = lhs_buffer.enqueue_fill(1.25)
    _ = rhs_buffer.enqueue_fill(2.5)

    # Wrap the device buffers in tensors
    var lhs_tensor = LayoutTensor[float_dtype, layout](lhs_buffer)
    var rhs_tensor = LayoutTensor[float_dtype, layout](rhs_buffer)
    var out_tensor = LayoutTensor[float_dtype, layout](out_buffer)

    # Calculate the number of blocks needed to cover the vector
    var grid_dim = ceildiv(VECTOR_WIDTH, BLOCK_SIZE)

    # Launch the vector_addition function as a GPU kernel
    ctx.enqueue_function_checked[vector_addition, vector_addition](
        lhs_tensor,
        rhs_tensor,
        out_tensor,
        VECTOR_WIDTH,
        grid_dim=grid_dim,
        block_dim=BLOCK_SIZE,
    )

    # Map to host so that values can be printed from the CPU
    with out_buffer.map_to_host() as host_buffer:
        var host_tensor = LayoutTensor[float_dtype, layout](host_buffer)
        print("Resulting vector:", host_tensor)
```

# GPU Programming in Mojo



- Unified programming for CPU + GPU
- The full power of standard CUDA/ROCm
  - Threads, warps, sync primitives
  - All the WMMA instructions
- "Without CUDA"
  - All GPU kernels written in Mojo.
  - Generate PTX directly without using CUDA toolkits or libraries
  - GPU kernels compiled on demand with small release container size
- GPU kernels for Nvidia, AMD, Apple in one language
- Library driven with "simple" compiler support

```

# =====
# thread_idx
# =====

@register_passable("trivial")
struct _ThreadIdx(Defaultable):
    """ThreadIdx provides static methods for getting the x/y/z coordinates of
    a thread within a block."""

    @always_inline("nodebug")
    fn __init__(out self):
        return

    @always_inline("nodebug")
    @staticmethod
    fn _get_intrinsic_name[dim: StringLiteral]() -> StaticString:
        @parameter
        if is_nvidia_gpu():
            return "llvm.nvvm.read.ptx.sreg.tid." + dim
        elif is_amd_gpu():
            return "llvm.amdgcn.workitem.id." + dim
        elif is_apple_gpu():
            return "llvm.air.thread_position_in_threadgroup." + dim
        else:
            return CompilationTarget.unsupported_target_error[
                StaticString,
                operation="thread_idx field access",
            ]()

    @always_inline("nodebug")
    fn __getattr__[dim: StringLiteral](self) -> UInt:
        """Gets the `x`, `y`, or `z` coordinates of a thread within a block.

        Returns:
            The `x`, `y`, or `z` coordinates of a thread within a block.
        """
        _verify_xyz[dim]()
        alias intrinsic_name = Self._get_intrinsic_name[dim]()
        return UInt(
            llvm_intrinsic[intrinsic_name, UInt32, has_side_effect=False]()
        )

alias thread_idx = _ThreadIdx()


```

**compile-time condition for different HW**

**llvm intrinsics**

# Recap: Mojo Compilation Pipeline

- Based on MLIR framework with LLVM as backend.
  - Parser
  - Parameter Domain passes (IR with generics/parameters)
  - Elaboration
  - Materialized IR optimizations (Parameters substituted with concrete values)
- Parallelized LLVM backend for **Fast** compilation time + **Performant** generated code.



What we've learned from building Mojo's optimization pipeline @ LLVM Dev 2024  
Parallelizing the LLVM Pipeline with MCLink @ EuroLLVM 2025

# Mojo GPU Compilation Flow

```
def main():
    constrained[has_accelerator(), "This example requires a supported GPU"]()

    # Get context for the attached GPU
    var ctx = DeviceContext()

    # Allocate data on the GPU address space
    var lhs_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)
    var rhs_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)
    var out_buffer = ctx.enqueue_create_buffer[float_dtype](VECTOR_WIDTH)

    # Fill in values across the entire width
    _ = lhs_buffer.enqueue_fill(1.25)
    _ = rhs_buffer.enqueue_fill(2.5)


    # Wrap the device buffers in tensors
    var lhs_tensor = LayoutTensor[float_dtype, layout](lhs_buffer)
    var rhs_tensor = LayoutTensor[float_dtype, layout](rhs_buffer)
    var out_tensor = LayoutTensor[float_dtype, layout](out_buffer)

    # Calculate the number of blocks needed to cover the vector
    var grid_dim = ceildiv(VECTOR_WIDTH, BLOCK_SIZE)

    # Launch the vector_addition function as a GPU kernel
    ctx.enqueue_function_checked[vector_addition, vector_addition](
        lhs_tensor,
        rhs_tensor,
        out_tensor,
        VECTOR_WIDTH,
        grid_dim=grid_dim,
        block_dim=BLOCK_SIZE,
    )


    # Map to host so that values can be printed from the CPU
    with out_buffer.map_to_host() as host_buffer:
        var host_tensor = LayoutTensor[float_dtype, layout](host_buffer)
        print("Resulting vector:", host_tensor)
```

GPU entry  
function as a  
parameter



# Mojo GPU Compilation Flow


Input Mojo program  
Host + GPU



```
module attributes {M.target_info = #M.target<triple = "arm64-apple-darwin23.6.0", arch = "apple-m2", features = "+aes,+bf16,+complxnum,+crc,+dotprod,+fp-armv8,+fp16ml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,+neon,+pauth,+perfmon,+ras,+rccpc,+rdm,+sha2,+sha3,+ssbs", data_layout = "e-m:o-p270:32:32-p271:32:32-p272:64:64-i128:128-n32:64-S128-Fn32", relocation_model = "pic", simd_bit_width = 128, index_bit_width = 64, kgen.env = #kgen.env<{}>} {  
    kgen.generator @hello() -> !kgen.none {  
        %none = kgen.param.constant none = <!kgen.none>  
        kgen.return %none : !kgen.none  
    }  
    kgen.generator export @entry(%arg0: !kgen.pointer<none>) {  
        kgen.param.declare nvptx: target = <!kgen.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">>  
        %0:2 = kgen.compile_offload<nvptx, 0, "", :() -> !kgen.none @hello -> !kgen.string, index  
        kgen.call @launch_kernel(%0#0, %0#1) : (!kgen.string, index) -> !kgen.none  
        kgen.return  
    }  
}
```

# Mojo GPU Compilation Flow

Input Mojo program  
Host + GPU




```
module attributes {M.target_info = #M.target<triple = "arm64-apple-darwin23.6.0", arch = "apple-m2", features = "+aes,+bf16,+complxnum,+crc,+dotprod,+fp-armv8,+fp16ml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,+neon,+pauth,+perfmon,+ras,+rccpc,+rdm,+sha2,+sha3,+ssbs", data_layout = "e-m:o-p270:32:32-p271:32:32-p272:64:64-i128:128-n32:64-S128-Fn32", relocation_model = "pic", simd_bit_width = 128, index_bit_width = 64, kgen.env = #kgen.env<{}>} {  
    kgen.generator @hello() -> !kgen.none {  
        %none = kgen.param.constant none = <#kgen.none>  
        kgen.return %none : !kgen.none  
    }  
    kgen.generator export @entry(%arg0: !kgen.pointer<none>) {  
        kgen.param.declare nvptx: target = <#kgen.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">>  
        %0:2 = kgen.compile_offload<nvptx, 0, "", :() -> !kgen.none @hello -> !kgen.string, index  
        kgen.call @launch_kernel(%0#0, %0#1) : (!kgen.string, index) -> !kgen.none  
        kgen.return  
    }  
}
```

```
module attributes {M.target_info = #M.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">, kgen.env = #kgen.env<{}>} {  
    kgen.generator export @hello() -> !kgen.none attributes {LLVMetadataArray = ["kgen.offload.kernelid", 0 : index]} {  
        %none = kgen.param.constant none = <#kgen.none>  
        kgen.return %none : !kgen.none  
    }  
}
```


# Mojo GPU Compilation Flow

Input Mojo program  
Host + GPU

\*.mojo (\*.🔥)



Slice out  
GPU Module



```

module attributes {M.target_info = #M.target<triple = "arm64-apple-darwin23.6.0", arch = "apple-m2", features = "+aes,+bf16,+complxnum,+crc,+dotprod,+fp-armv8,+fp16ml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,+neon,+pauth,+perfmon,+ras,+rccpc,+rdm,+sha2,+sha3,+ssbs", data_layout = "e-m:o-p270:32:32-p271:32:32-p272:64:64-i128:128-n32:64-S128-Fn32", relocation_model = "pic", simd_bit_width = 128, index_bit_width = 64>, kgen.env = #kgen.env<{}>} {
    kgen.generator @hello() -> !kgen.none {
        %none = kgen.param.constant none = <!kgen.none>
        kgen.return %none : !kgen.none
    }
    kgen.generator export @entry(%arg0: !kgen.pointer<none>) {
        kgen.param.declare nvptx: target = <#kgen.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">>
        %0:2 = kgen.compile_offload<nvptx, 0, "", :() -> !kgen.none @hello -> !kgen.string, index
        kgen.call @launch_kernel(%0#0, %0#1) : (!kgen.string, index) -> !kgen.none
        kgen.return
    }
}
  
```

```

module attributes {M.target_info = #M.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">, kgen.env = #kgen.env<{}>} {
    kgen.generator export @hello() -> !kgen.none attributes {LLVMetadataArray = ["kgen.offload.kernelid", 0 : index]} {
        %none = kgen.param.constant none = <!kgen.none>
        kgen.return %none : !kgen.none
    }
}
  
```

PTX output

```


"//\0A// Generated by LLVM NVPTX Back-End\0A//\0A\0A.version 7.0\0A.target sm_80\0A.address_size 64\0A\0A\09// .
globl\09hello\0A\0A.visible .entry hello()\0A{\0A\0A\0A\09ret;\0A\0A}\0A"
  
```

# Mojo GPU Compilation Flow


Input Mojo program  
Host + GPU

\*.mojo (\*.🔥)

\*.mojom pkg (+.📦)



Slice out  
GPU Module



PTX output

```
// Generated by LLVM NVPTX Back-End\0A//\0A\0A.version 7.0\0A.target sm_80\0A.address_size 64\0A\0A\09// .  
globl\09hello\0A\0A.visible .entry hello()\0A{\0A\0A\0A\09ret;\0A\0A}\0A"
```

Host code with  
compiled offload PTX

```

module attributes {M.target_info = #M.target<triple = "arm64-apple-darwin23.6.0", arch = "apple-m2", features = "+aes,+bf16,+complxnum,+crc,+dotprod,+fp-armv8,+fp16fml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,+neon,+pauth,+perfmon,+ras,+rcpc,+rdm,+sha2,+sha3,+ssbs", data_layout = "e-m:o-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-n32:64-S128-Fn32", relocation_model = "pic", simd_bit_width = 128, index_bit_width = 64>, kgen.env = #kgen.env<{}>} {
  kgen.generator @hello() -> !kgen.none {
    %none = kgen.param.constant none = <#kgen.none>
    kgen.return %none : !kgen.none
  }
  kgen.generator export @entry(%arg0: !kgen.pointer<none>) {
    kgen.param.declare nvptx: target = <#kgen.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">, kgen.env = #kgen.env<{}>> {
      %0:2 = kgen.compile_offload<nvptx, 0, "", :() -> !kgen.none @hello -> !kgen.string, index
      kgen.call @launch_kernel(%0#0, %0#1) : (!kgen.string, index) -> !kgen.none
      kgen.return
    }
  }
}
  
```

```


module attributes {M.target_info = #M.target<triple = "nvptx64-nvidia-cuda", arch = "sm_80", simd_bit_width = 128, index_bit_width = 64, tune_cpu = "sm_80">, kgen.env = #kgen.env<{}>} {
  kgen.generator export @hello() -> !kgen.none attributes {LLVMMetadataArray = ["kgen.offload.kernelid", 0 : index]} {
    %none = kgen.param.constant none = <#kgen.none>
    kgen.return %none : !kgen.none
  }
}
  
```

```

module attributes {M.target_info = #M.target<triple = "arm64-apple-darwin23.6.0", arch = "apple-m2", features = "+aes,+bf16,+complxnum,+crc,+dotprod,+fp-armv8,+fp16fml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,+neon,+pauth,+perfmon,+ras,+rcpc,+rdm,+sha2,+sha3,+ssbs", data_layout = "e-m:o-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-n32:64-S128-Fn32", relocation_model = "pic", simd_bit_width = 128, index_bit_width = 64>, kgen.env = #kgen.env<{}>} {
  kgen.func export @entry(%arg0: !kgen.pointer<none>) {
    %string = kgen.param.constant string = <"/\0A// Generated by LLVM NVPTX Back-End\0A//\0A\0A.version 7.0\0A.target sm_80\0A.address_size 64\0A\0A\09// .  
globl\09hello\0A\0A.visible .entry hello()\0A{\0A\0A\0A\09ret;\0A\0A}\0A">
    %index0 = kgen.param.constant = <0>
    kgen.call @launch_kernel(%string, %index0) : (!kgen.string, index) -> !kgen.none
    kgen.return
  }
}
  
```

# Compiler support for Debugging Kernels

- Inspecting Mojo GPU kernel in LLVM IR, assembly (PTX), or object file (cubin, hsaco)
- Same GPU compilation flow, just change what the pipeline produces.
- Plug the result as string to host code.



**asm**

```

// Generated by LLVM NVPTX Back-End
//
.version 8.1
.target sm_80
.address_size 64

// .globl compile_offload_hello6A6AoApA

.visible .entry compile_offload_hello6A6AoApA()
{
    ret;
}

```

**Ilvm IR**

```

; ModuleID = 'compile_offload.mojo'
source_filename = "compile_offload.mojo"
target datalayout =
"e-p3:32:32-p4:32:32-p5:32:32-p6:32:32-p7:32:32-i64:64-i128:128-i256:256-v16:16-v32:32-n16:32:64"
target triple = "nvptx64-nvidia-cuda"

; Function Attrs: norecurse
define dso_local ptx_kernel void @compile_offload_hello6A6AoApA() #0 {
    ret void
}

attributes #0 = { norecurse "target-cpu"="sm_80" "target-features"="+ptx81,+sm_80" "tune-cpu"="sm_80" }

!llvm.module.flags = !{!0}

!0 = !{i32 2, !"Debug Info Version", i32 3}

```

**opt Ilvm IR**

```

; ModuleID = 'compile_offload.mojo'
source_filename = "compile_offload.mojo"
target datalayout = "e-m:0-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-n32:64-S128-Fn32"
target triple = "arm64-apple-darwin23.6.0"

; Function Attrs: mustprogressnofree norecurse nosync nounwind willreturn memory(none) "target-cpu"="apple-m2"
"target-features"="+aes,+bf16,+comp_lxnum,+crc,+dotprod,+fp-armv8,+fp16fml,+fpac,+fullfp16,+i8mm,+jsconv,+lse,
+neon,+pauth,+perfmmon,+ras,+rcpc,+rdm,+sha2,+sha3,+ssbs" )

!llvm.module.flags = !{!0}

!0 = !{i32 2, !"Debug Info Version", i32 3}

```

Compiler support for Debugging Kernels

```

@fieldwise_init
@register_passable("trivial")
struct _Info:
    var kernel: __mlir_type.`!kgen.string`
    var name: __mlir_type.`!kgen.string`
    var num_captures: __mlir_type.index

@fieldwise_init
@register_passable("trivial")
struct Info:
    var kernel: StaticString
    var name: StaticString
    var num_captures: Int

@always_inline
fn _compile_info[
    func_type: AnyTrivialRegType, //,
    func: func_type,
    ,
    emission_kind: StaticString = "asm",
    target: _TargetType = A100.target(),
    compile_options: StaticString = CompilationTarget[
        target
    ].default_compile_options(),
]() -> Info:
    var info = __mlir_op.`kgen.compile_offload`[
        target_type=target,
        emission_kind = _get_emission_kind_id[emission_kind]().__mlir_value,
        emission_option = _get_kgen_string[compile_options](),
        func=func,
        _type=_Info,
    ]()

    return Info(
        StaticString(info.kernel),
        StaticString(info.name),
        Int(__mlir_value=info.num_captures),
    )

fn hello():
    pass

fn main():
    # compile kernel to asm.
    t1 = _compile_info[hello, emission_kind="asm"]()
    print(t1.kernel)

    # compile kernel to llvm ir.
    t2 = _compile_info[hello, emission_kind="llvm"]()
    print(t2.kernel)


    # compile kernel to optimized llvm ir.
    t3 = _compile_info[hello, emission_kind="llvm-opt"]()
    print(t3.kernel)

```

# Per-kernel GPU compilation




- Multiple GPU kernels in one Mojo program or an MAX Model.
  - Slice each kernel out to run full GPU pipeline.
  - Ease to debug each kernel for lowering from input MLIR to LLVM.
  - Efficient to cache compiled kernels if they will be used in another program or model.
  - Logically easy to maintain.



# All-kernel GPU compilation



- Multiple GPU kernels in one Mojo program or an MAX Model.
  - Slice out one GPU model with all the GPU kernels.
  - Only run MLIR pipeline once instead of per kernel.
  - Split each kernel into separate LLVM modules to run backend pipeline in parallel. \*
  - Can cache LLVM compilation for kernels if they will be used in another program or model.
  - Faster compilation time.



\* *Parallelizing the LLVM Pipeline with MCLink* @ EuroLLVM 2025

# Compilation Time Comparison

/usr/bin/time mojo build Kernels/test/gpu/linalg/test\_matmul.mojo

g5.8xlarge

opt, dbg-level	All-kernel GPU compilation flow					Per-kernel GPU compilation flow					user+sys old/new x	maxres new/old x
	user (s)	system (s)	elapsed	CPU	maxresident(k)	user (s)	system (s)	elapsed	CPU	maxresident(k)		
O3, none	236.89	9.40	0:51.07	482%	3138480	499.05	18.17	0:56.40	917%	4540640	2.10	0.691197716621445
O3, line-tables	292.20	35.73	1:08.71	477%	36492336	590.43	22.34	1:05.53	934%	5295764	1.87	6.89085389756794
O2, none	259.18	10.59	0:52.61	512%	3604328	585.77	21.04	1:02.15	976%	5484716	2.25	0.657158547498175
O2, line-tables	319.19	37.70	1:11.42	499%	38531356	679.10	25.24	1:11.65	983%	6226044	1.97	6.18873814576318

/usr/bin/time mojo build Kernels/test/gpu/linalg/test\_linalg\_matmul\_gpu.mojo

g5.8xlarge

opt, dbg-level	All-kernel GPU compilation flow					Per-kernel GPU compilation flow					usr+sys old/new x	maxres new/old x
	user (s)	system (s)	elapsed	CPU	maxresident(k)	user (s)	system (s)	elapsed	CPU	maxresident(k)		
O3, none	49.44	2.03	0:20.95	245%	833268	60.36	3.80	0:20.67	310%	1775324	1.25	0.469361085638452
O3, line-tables	57.53	3.00	0:23.99	252%	1450892	70.53	4.16	0:23.45	318%	2045704	1.23	0.709238482204659
O2, none	52.37	2.29	0:21.30	256%	919580	66.36	3.94	0:21.27	330%	1998244	1.29	0.46019405037623
O2, line-tables	59.88	3.32	0:23.95	263%	1546068	75.69	4.75	0:23.45	342%	2238268	1.27	0.690743020942979

MODULAR\_MAX\_TEMPS\_DIR=/tmp/llama3\_artifacts flowmeter -n 1 utils/benchmarking/flowmeter/pipelines/max-llama3\_1-python-gpu.yaml

a100

	input_toks	output_toks	startup_ms	TTFT_ms	CE_tps	Time per Output Token	TG_tps	Total Latency	Total Throughput
All-kernel	482	128	61263.26	2738.74	175.99	13.06	76.57	4397.29	0.23
Per-kernel	482	128	227997.04	3376.99	142.73	13.03	76.74	5031.90	0.20
Per-kernel/ All-kernel(x)			3.72	1.23	0.81	1.00	1.00	1.14	0.87

# Conclusions

- Mojo provides a unified way to write CPU+GPU code => **one MLIR module**
- Library driven GPU feature implementation for different vendors => **simple compiler, no heroic magic**
- MLIR unlocks seamless compiler integration:
  - Native compiler support to know what to slice into GPU mlir modules during elaboration.
  - Add an MLIR op *kgen.compile\_offload*.
  - Mojo has support as MLIR sugar:
    - No need to change the parser to extend syntax.
    - Ease of writing library code that uses architecture specific dialects (NVVM, ROCDL) and low-level intrinsics (LLVM).
- Compiling multiple kernels in one MLIR module and split for LLVM pipeline => **fast compilation**
- **Generally applicable** to compile other accelerator offloads built on top of LLVM/MLIR framework.
- Mojo GPU kernels are all open-sourced

<https://github.com/modular/modular/tree/main/max/kernels>



# Thank you!

Acknowledgement:

The Mojo Language Team at Modular

