

Modular



# Mojo + PyTorch: A Simpler Path to Custom Kernels

Spenser Bauman

Oct 22, 2025

Compiler Engineer @  
MODULAR INC.

# Agenda

- 
- 01 Why Mojo & MAX?
  - 02 Write a custom kernel in Mojo
  - 03 Interface to it easily from PyTorch
  - 04 Good performance without the pain
- 

✖What you won't see:

- Hardware lock-in
- PyBind, CMake, setuptools
- Performance gaps
- CUDA version mismatches
- Library path problems



# Why Mojo?



# A look at existing performance libraries



# Many are written in assembly...

... so much for abstraction  
even across members of an  
architectural family ...

```
lea      rax,[rdx+r8*2]
vpmovzxbw ymm4,XMMWORD PTR [rdx]
vpmovzxbw ymm5,XMMWORD PTR [rdx+r8]
vpmovzxbw ymm6,XMMWORD PTR [rax]
vpmovzxbw ymm7,XMMWORD PTR [rax+r8]
lea      rax,[rcx+r11*4]
vmovdqu YMMWORD PTR [rcx],ymm4
vmovdqu YMMWORD PTR [rcx+r11*2],ymm5
vmovdqu YMMWORD PTR [rax],ymm6
vmovdqu YMMWORD PTR [rax+r11*2],ymm7
vpaddw  ymm0,ymm0,ymm4
vpaddw  ymm1,ymm1,ymm5
vpaddw  ymm2,ymm2,ymm6
vpaddw  ymm3,ymm3,ymm7
add     rdx,16
add     rcx,16*2
sub    rbx,16
```

# C++ Templates



Source: Composable Kernels

```
static constexpr auto GemmDefault =
    ck::tensor_operation::device::GemmSpecialization::Default;

using DeviceGemmInstance = ck::tensor_operation::device::DeviceGemmXdl<
    ADataType, BDataType, CDataType, AccDataType, ALayout, BLayout, CLayout,
    AELEMENTOp, BELEMENTOp, CELEMENTOp, GemmDefault, 256, 128, 128, 4, 2, 16,
    16, 4, 4, S<4, 64, 1>, S<1, 0, 2>, S<1, 0, 2>, 2, 2, 2, true, S<4, 64, 1>,
    S<1, 0, 2>, S<1, 0, 2>, 2, 2, 2, true, 7, 1>;

using ReferenceGemmInstance =
    ck::tensor_operation::host::ReferenceGemm<ADatatype, Bdatatype, Cdatatype,
    Accdatatype, AELEMENTOp,
    BELEMENTOp, CELEMENTOp>;

#include "run_gemm_example.inc"
```

# C++ DSL for ASM



Source: OneDNN

```
L(labels[4]);
test(K, 2);
jle(labels[5], T_NEAR);
innerkernel2(unroll_m, unroll_n, isLoad1Unmasked, isLoad2Unmasked, isDirect,
            |           |           |           |
            |           |           |           isCopy, useFma, reg00, reg01, reg02, reg03, reg04, reg05,
            |           |           |           reg06, reg07, reg08, reg09, reg10, reg11, reg12, reg13, reg14,
            |           |           |           reg15, reg16, reg17, reg18, reg19, reg20, reg21, reg22, reg23);
align(16);

L(labels[5]);
if (unroll_m == 16) {
    if (unroll_n <= 3) {
        vaddps(reg00, reg00, reg12);
        vaddps(reg01, reg01, reg13);
        vaddps(reg02, reg02, reg14);
        vaddps(reg06, reg06, reg18);
        vaddps(reg07, reg07, reg19);
        vaddps(reg08, reg08, reg20);
    }
}
```

# Python program to generate ASM



Source: Tensile

```
for iui in range(0, innerUnroll):
    for idx1 in range(0, kernel["ThreadTile1"]):
        for idx0 in range(0, kernel["ThreadTile0"]):
            vars["idx0"] = idx0
            vars["idx1"] = idx1
            vars["a"] = idx0 if writer.tPB["tile01Idx"] else idx1
            vars["b"] = idx1 if writer.tPB["tile01Idx"] else idx0
            vars["iui"] = iui

            vars["cStr"] = "v[vgprValuC + {idx0} + {idx1}*{ThreadTile0}].format_map(vars)
            vars["aStr"] = "v[vgprValuA_X{m}_I{iui} + {a}].format_map(vars)
            vars["bStr"] = "v[vgprValuB_X{m}_I{iui} + {b}].format_map(vars)

            if instruction == "v_fma_f32":
                kStr += "v_fma_f32 {cStr}, {aStr}, {bStr}, {cStr}{endLine}.format_map(vars)
            else:
                kStr += "{instruction} {cStr}, {aStr}, {bStr}{endLine}.format_map(vars)

            kStr += priority(writer, 1, "Raise priority while processing macs")
```

# Template (*but not templates*) to generate C++

Source: XNNPack

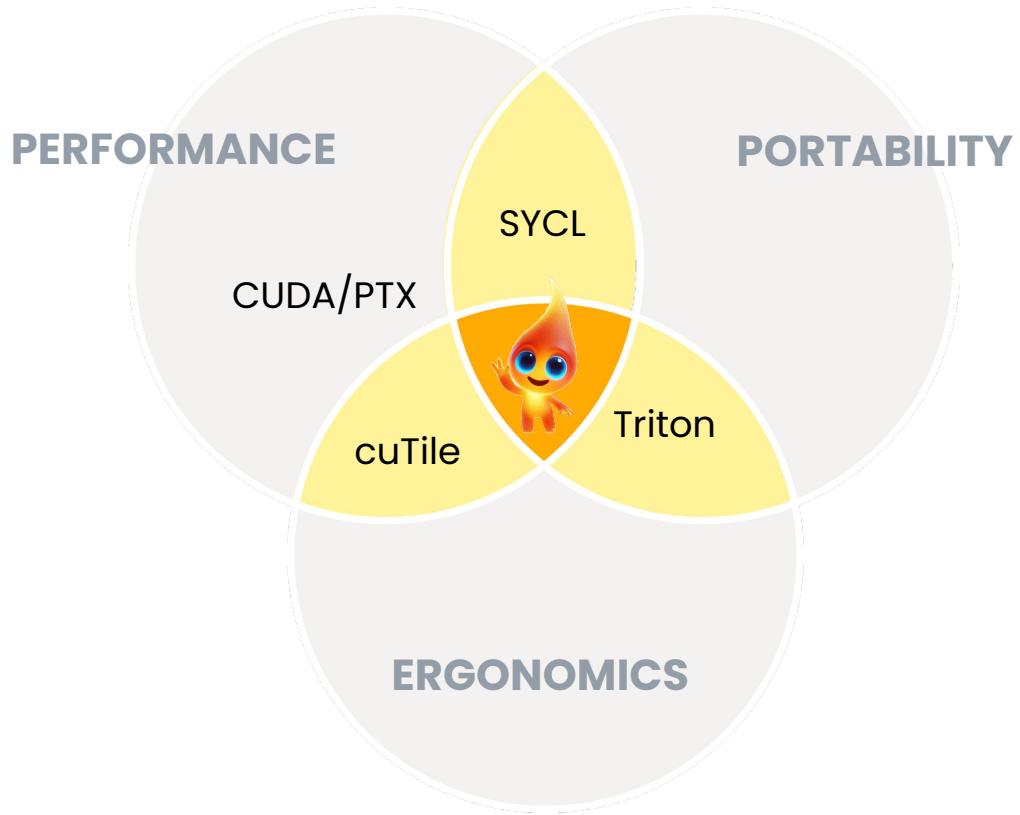
```
const __m128i vsign_mask =
    _mm_load_si128((const __m128i*)params->${PARAMS_STRUCT}.sign_mask);
const __m256 vsat_cutoff = _mm256_load_ps(params->${PARAMS_STRUCT}.sat_cutoff);
const __m256 vlog2e = _mm256_load_ps(params->${PARAMS_STRUCT}.log2e);
const __m256 vmagic_bias = _mm256_load_ps(params->${PARAMS_STRUCT}.magic_bias);
const __m256 vminus_ln2 = _mm256_load_ps(params->${PARAMS_STRUCT}.minus_ln2);
$for i in reversed(range(2, P + 1))
    : const __m256 vc${i} = _mm256_load_ps(params->${PARAMS_STRUCT}.c${i});
$if P != H + 1 : const __m256 vminus_one =
    _mm256_load_ps(params->${PARAMS_STRUCT}.minus_one);
const __m256 vtwo = _mm256_load_ps(params->${PARAMS_STRUCT}.two);
$if P == H + 1 : const __m256 vminus_one =
    _mm256_load_ps(params->${PARAMS_STRUCT}.minus_one);
```

# Pervasive suffering

Maintainability,  
debugging, tooling, ...



**Choose:  
Performance or  
Ease of Use**



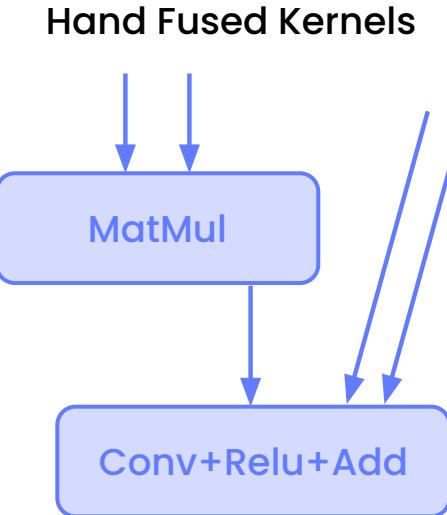
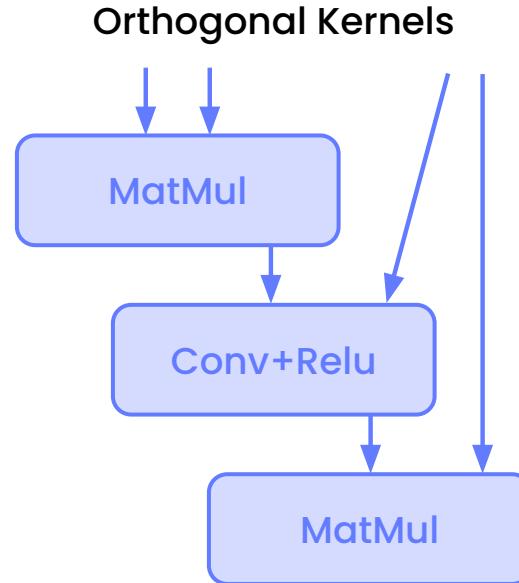
---

Why MAX?

---

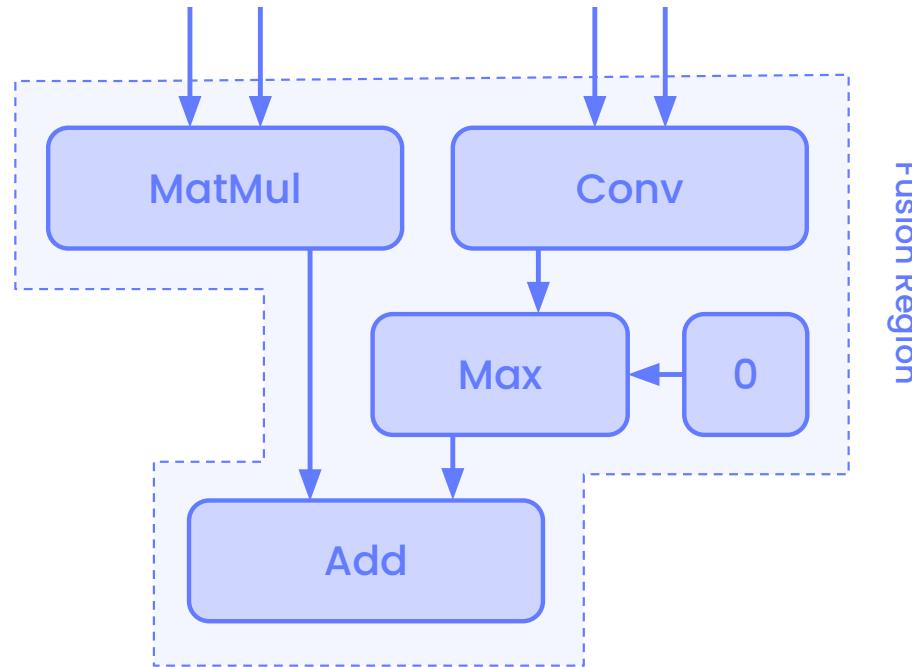


# It started with Frameworks with Large Kernel Libraries



Thousands of kernels later: it doesn't scale!

# Then: AI Compilers To Rescue?



There's a problem ... generality!

# But: Where is Generality?

## Many common limitations...

- Dynamic shapes
- Sparsity
- Quantization
- Custom ops
- Embedded support
- Model coverage

## GenAI changed the game:

- Continuous innovation in attention and other mechanisms: more explosion of fused kernels
- Diversity of KVCache optimizations
- Performance/TCO is critical for inference

*"Generality is, indeed, an indispensable ingredient of reality; for mere individual existence or actuality without any regularity whatever is a nullity. Chaos is pure nothing.*

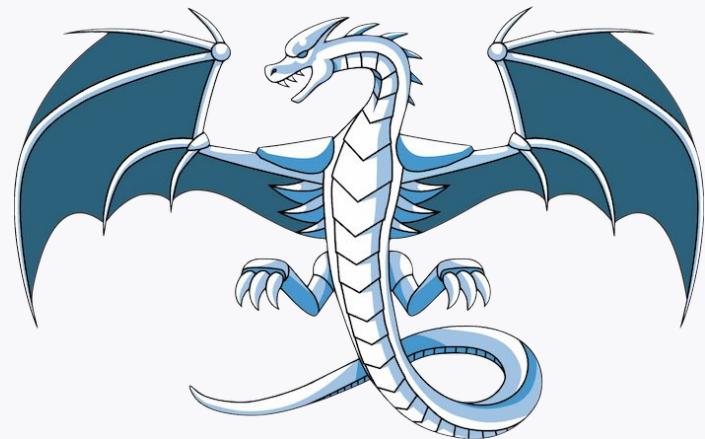
- Charles Sanders Peirce

# Challenge: Compiler Scalability!

**Difficult to hire compiler engineers ...**

- ... who have AI modeling experience, and
- ... who know exotic numerics, and
- ... who know specialized HW details

**AI Research cannot rely on:  
"compiler engineer in-the-loop"!**





# To the rescue

## MAX – Modular Accelerated eXecution

- **Mojo Programming Language**: System-level performance with Python syntax
- **Python Graph API**: Familiar PyTorch-like interface
- **Graph Compiler**: Automatic optimization and fusion
- **Optimized GPU Kernels**: Hardware-specific optimizations
- **Hardware Abstraction**: Write once, run anywhere

# Why



## A Pythonic System Programming Language

- Modern high-level systems language
- Hardware agnostic (CPU and GPU)
- Memory safety, powerful metaprogramming system
- Full toolchain + VS Code + LSP support
- Best of both worlds: productivity AND performance

## Cross-Platform Hardware Support

- Supports CPU and GPU targets without vendor lock-in
- Works with Nvidia, AMD (Apple Silicon in nightlies)
- Significantly smaller container sizes vs competitors (600MB vs 11GB)

Tag summary

Recent tags: **nightly**

Content type: Image

Digest: sha256:db5ee2eb8...

**Size: 581.7 MB**

Last updated: about 6 hours ago

```
docker pull modular/max-nvidia-base:nightly
```

Tag summary

Recent tags: **nightly**

Content type: Image

Digest: sha256:a8419d767...

**Size: 11.5 GB**

Last updated: about 14 hours ago

```
docker pull vllm/vllm-openai:nightly
```

---

# Write a custom kernel in Mojo

---



# Anatomy of a Mojo Kernel

Fusion  
Properties

Parallelize/Vectorize  
across CPU or GPU

```
@compiler.register("grayscale")
struct Grayscale:
    @staticmethod
    fn execute[ target: StaticString,
    ](
        img_out: FusedOutputTensor[dtype=float32, rank=2],
        img_in: InputTensor[dtype=float32, rank=3],
        ctx: DeviceContextPtr,
    ) raises:
        @parameter
        fn color_to_grayscale[
            width: Int
        ](idx: IndexList[img_out.rank]) -> SIMD[float32, width]:
            var row, col = idx[0], idx[1]
            var r = img_in.load[width](Index(0, row, col))
            var g = img_in.load[width](Index(1, row, col))
            var b = img_in.load[width](Index(2, row, col))
            return 0.21 * r + 0.71 * g + 0.07 * b
    foreach[color_to_grayscale, target=target](img_out, ctx)
```

Registered name

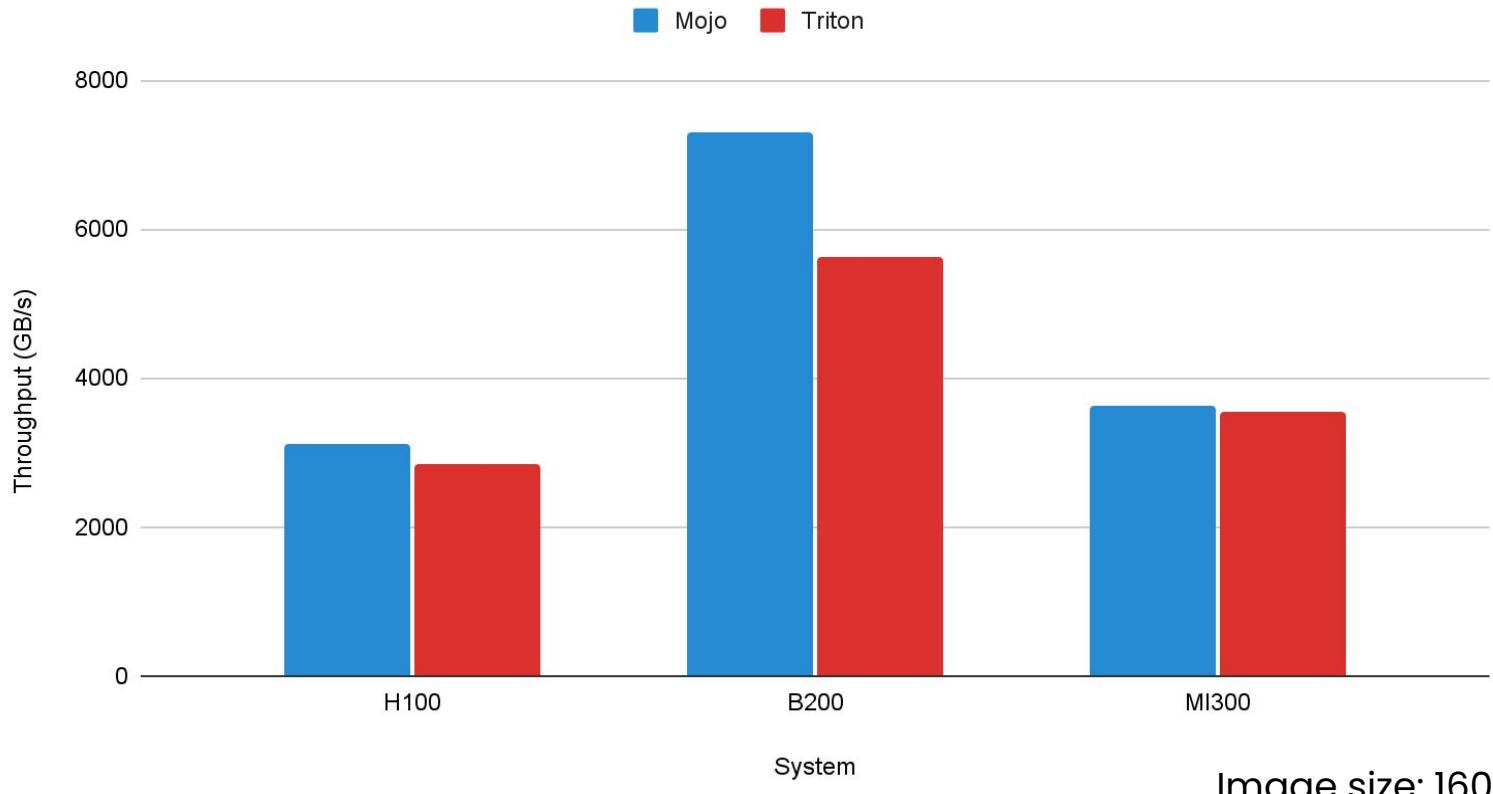
Entry point

Structured Inputs  
(not pointers)

Modular

# And how does it perform?

Triton vs Mojo Grayscale Throughput



# What about something more realistic?

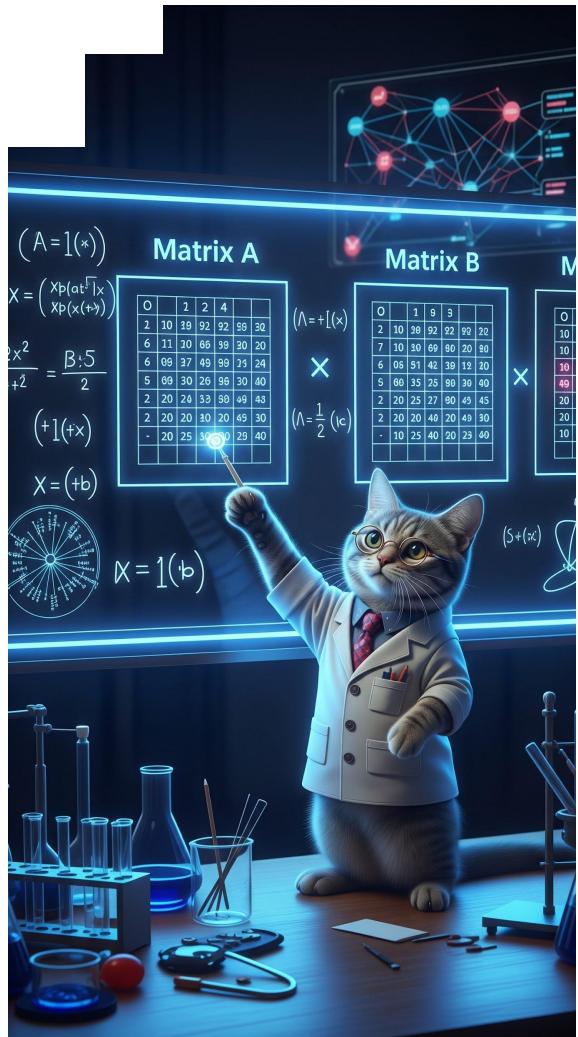
At Modular, all our kernels are implemented in Mojo

👤 ... by a relatively small team of engineers (~10 people),

🔧 ... for all the hardwares we support: CPUs + 7 major GPU architectures from 3 different vendors (NV, AMD, Apple),

🚀 ... and get SOTA throughput on production models,

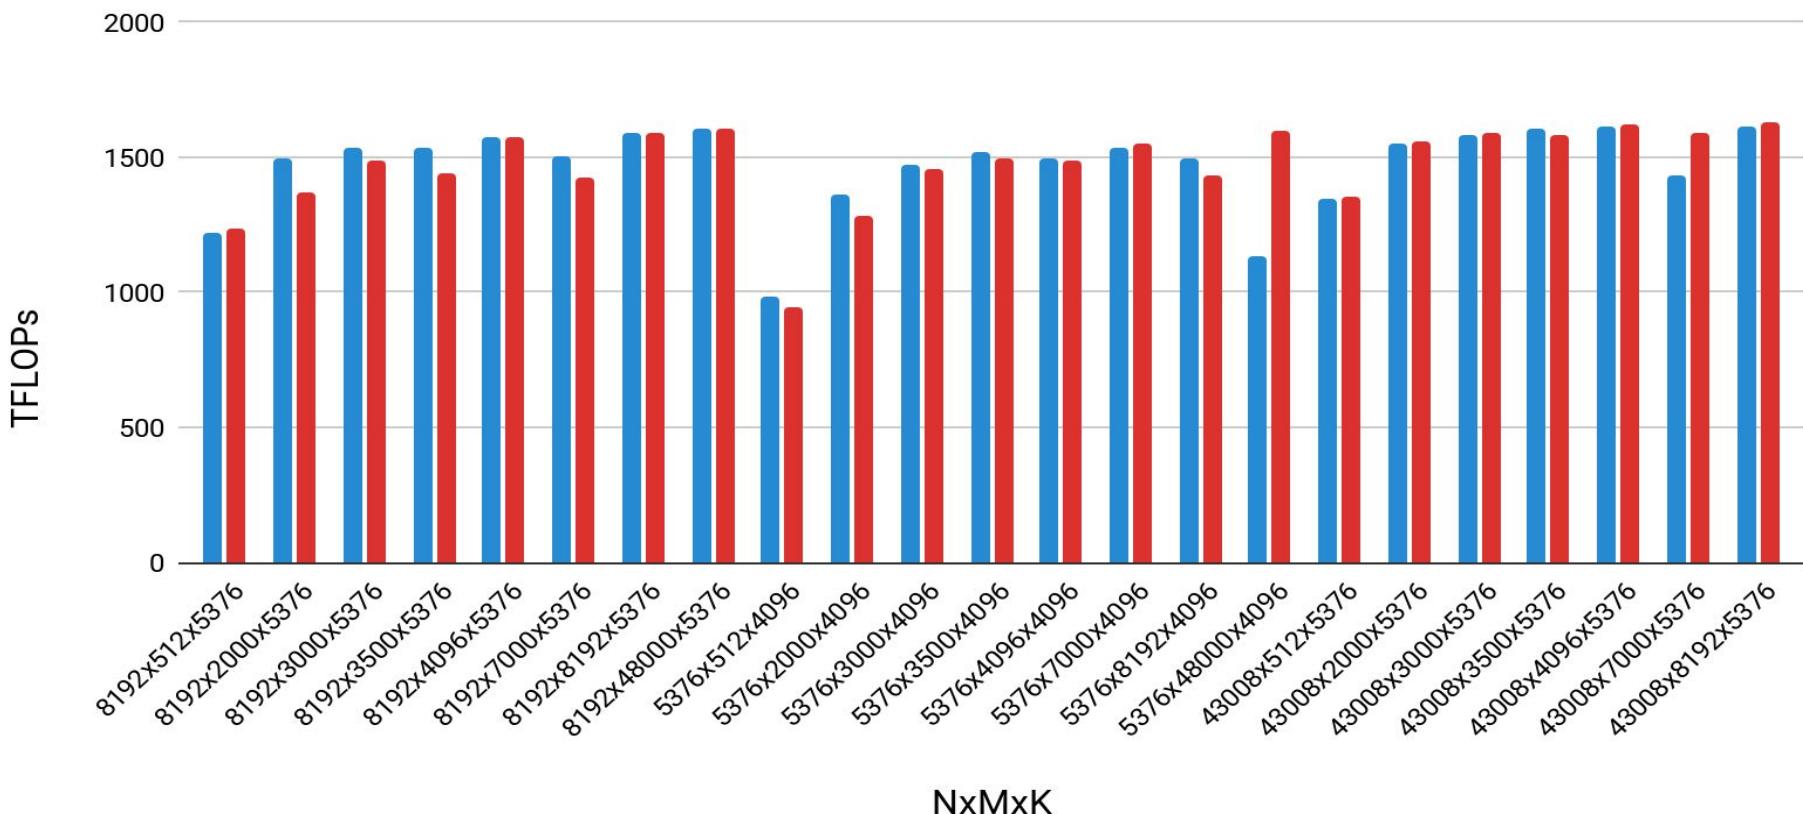
🔥 ... so how do we do on performance critical kernels?



Mojo vs cuBLAS (B200)

# Matrix Multiplication

Mojo    cuBLAS



# 2 APIs to Integrate Mojo into PyTorch

1.  An API to directly call custom kernels
2.  An API to call full MAX graphs



Modular

# Call Your Kernel from PyTorch

Load your library of custom kernels (no compile step):

```
custom_ops = CustomOpLibrary("path/to/ops")
```



Wrap the destination passing kernel:

```
@torch.compile
def max_grayscale(x: torch.Tensor) -> torch.Tensor:
    result = torch.empty(x.shape[1:], dtype=x.dtype, device=x.device)
    custom_ops.grayscale(result, x)
    return result
```



Compatible with “fullgraph=True”



# When to use this API

## Tradeoffs for this API:

- ✓ Simple for a small number of kernels
- ✗ Overhead due to \_\_dpack\_\_ & stream synchronization
- ✗ Introduced optimization boundaries (no fusion)

# Integrate with MAX Graphs



## This API:

- ✗ More work for single ops
- ✓ Scales beyond a single operation
- ✓ Less marshalling/dispatch overhead
- ✓ More optimization opportunities

```
@max.torch.graph_op
def max_grayscale_graph(
    pic: max.graph.TensorValue
) -> max.graph.TensorValue:
    c, h, w = pic.shape

    scale = ops.constant(
        np.array([0.21, 0.71, 0.07]),
        dtype=DType.float32,
        device=DeviceRef.GPU()
    ).reshape([c, 1, 1])

    scaled = pic * ops.broadcast_to(scale, (c, h, w))
    grayscaled = ops.sum(scaled, axis=0)
    # max reductions don't remove the dimension, need to squeeze
    return ops.squeeze(grayscaled, axis=0)
```

# Key Takeaways



1. Prototype and optimize kernels in one system
2. No vendor lock-in
3. NEW Easy integration with PyTorch





# Thank You

Custom ops tutorial  
[modul.ar/ptc-ops](https://modul.ar/ptc-ops)

Support & discussion  
[forum.modular.com](https://forum.modular.com)

Optimizing Matmul on Blackwell:  
<https://www.modular.com/matrix-multiplication-on-blackwell>

